



BIRLA INSTITUTE OF TECHNOLOGY & SCIENCE, PILANI

WORK INTEGRATED LEARNING PROGRAMMES

COURSE HANDOUT

Part A: Content Design

Course Title	Big Data Systems
Course No(s)	DSECL ZG522
Credit Units	5
Course Author	Prof. Shan Balasubramaniam
Version No	2.0
Last Revised By	Pravin Y Pawar
Date	1 January 2020

Course Description

The course introduces the students to the concepts of Systems for Analytics with particular emphasis on processing Big Data. It introduces distributed computing models for storage and processing of Big Data with specific coverage of block storage, file systems, and databases on the one hand and batch processing, in-memory distributed processing, and stream processing on the other. Hadoop (along with associated technologies such as Hive and Pig), Spark, and Amazon's storage and database services are used as exemplar platforms.

Course Objectives

CO1	Enable students to understand requirements for and constraints in storing and processing Big Data
CO2	Enable students to leverage commodity infrastructure (such as scale-out clusters, distributed data stores, and the cloud) and the appropriate platforms and services for storing and processing Big Data.
CO3	Enable students to implement solutions for big data processing
CO4	Enable students to develop a working knowledge of stream processing

Text Book(s)

T1	Seema Acharya and Subhashini Chellappan. <i>Big Data and Analytics</i> . Wiley India Pvt. Ltd. Second Edition
----	---

Reference Book(s) & other resources

R1	DT Editorial Services. <i>Big Data - Black Book</i> . DreamTech. Press. 2016
R2	Kai Hwang, Jack Dongarra, and Geoffrey C. Fox. <i>Distributed and Cloud Computing: From Parallel Processing to the Internet of Things</i> . Morgan Kauffman 2011
R3	<i>Additional Reading</i> (as assigned for specific topics)

Learning Outcomes:

No	Learning Outcomes
LO1	A comprehensive understanding of the Big Data ecosystem and along with the typical technologies involved.
LO2	Apply concepts from distributed computing and use the Hadoop/Map-reduce framework and for solving typical big data problems.
LO3	Identify and use appropriate storage / database platforms for Big data storage along with appropriate querying mechanisms / interfaces for retrieval.
LO4	Use in-memory processing and stream processing techniques for building Big Data systems.

Modular Structure

Module #	Name of Module	Contact Sessions
1	Data Engineering	1 - 2
2	Big Data Analytics	3 - 5
3	Hadoop Ecosystem	6 - 9
4	Big Data Storages	10 - 12
5	Spark for Big Data Processing	13 - 16

Part B: Contact Session Plan

Academic Term	II Semester 2019-2010
Course Title	Big Data Systems
Course No	DSECL ZG522
Lead Instructor	Pravin Y Pawar

Session #	Contact Hours(#)	List of Topic / Title	Text/Ref Book/external resource
1	1	Different Types of Data and Storage for Data: Structured Data (Relational Databases) , Semi-structured data (Object Stores), and Unstructured Data (File systems) What is Big Data? Characteristics of Big Data.	T1 Ch. 1 T1 Ch.2

		Systems perspective - Processing: In-memory vs. (from) secondary storage vs. (over the) network	R2 Sec 1.2.3
	2	Locality of Reference: Principle, examples Impact of Latency: Algorithms and data structures that leverage locality, data organization on disk for better locality	Class Slides
2	3	Parallel and Distributed Processing: Motivation (Size of data and complexity of processing); Storing data in parallel and distributed systems: Shared Memory vs. Message Passing; Strategies for data access: Partition, Replication, and Messaging.	R2 Sec. 1.2, 1.3.4, and 1.4.1
	4	Memory Hierarchy in Distributed Systems: In-node vs. over the network latencies, Locality, Communication Cost. Distributed Systems: Motivation (size, scalability, cost-benefit), Client-Server vs. Peer-to-Peer models, Cluster Computing: Components and Architecture	Class Slides R2 Sec. 2.1 to 2.3
3	5	Big Data Analytics: Requirements, constraints, approaches, and technologies.	T1 Sec. 3.1 to 3.11; R1 Ch. 3 and Ch. 6
	6	Big Data Systems – Characteristics: Failures; Reliability and Availability; Consistency – Notions of Consistency.	T1 Ch. 4 AR
4	7	CAP Theorem and implications for Big data Analytics	T1 Sec. 3.12 and 3.13; AR
	8	Big Data Lifecycle: Data Acquisition, Data Extraction – Validation and Cleaning, Data Loading, Data Transformation, Data Analysis and Visualization. Case study – Big data application	T1 Sec. 2.9 to 2.12; R1 Ch. 6 and Ch. 7
5	9-10	Distributed Computing - Design Strategy: Divide-and-conquer for Parallel / Distributed Systems - Basic scenarios and Implications. Programming Patterns: Data-parallel programs and <i>map</i> as a construct; Tree-parallelism, and <i>reduce</i> as a construct; Map-reduce model: Examples (of map, reduce, map-reduce combinations, and Iterative map-reduce)	AR
6	11-12	Hadoop: Introduction, Architecture, and Map-reduce Programming on Hadoop	T1 Sec. 5.1 and 5.2, Sec. 5.7, Sec. 5.11, and Ch. 8; R1 Ch. 5 and Ch. 9; R2 Sec. 1.4.3 and 6.2.2; AR
7	13-14	Hadoop: Hadoop Distributed File System (HDFS),	T1 5.10 and 5.12;

		Scheduling in Hadoop (using YARN). Example – Hadoop application.	R1 Ch. 4 (sections on HDFS and Yarn) and Ch. 11; AR
8	15-16	Hadoop Ecosystem: Databases and Querying (HBASE, Pig, and Hive)	T1 Sec. 5.13; R1 Ch. 4 (sections on HBase, Hive, and Pig) and Ch. 5 (section on HBase)
9	17-18	Hadoop Ecosystem: Integration and coordination (Sqoop, Flume, Zookeeper & Oozie)	T1 Sec. 5.13; R1 Ch. 4 (sections on Sqoop, Flume, Zookeeper & Oozie)
10	19-20	NoSQL databases: Introduction, Architecture, Querying, Variants, Case Study.	T1 Sec. 4.1, Ch. 6, and Ch. 7
11	21	Cloud Computing: A brief overview: Motivation, Structure and Components; Characteristics and advantages – Elasticity. Services on the cloud.	AR
	22	Storage as a Service: Forms of storage on the cloud, databases on the cloud.	AR
12	23	Amazon's storage services: block storage, file system, and database; EBS, SimpleDB, S3	AR (<i>sourced from Amazon</i>)
	24	Case study – Amazon DynamoDB (Access/Querying model, Database architecture and applications on the cloud).	-
13	25	Spark: Introduction, Architecture and Features	AR
	26	Programming on Spark: Resilient Distributed Datasets, Transformation, Examples	AR (Apache Spark docs.)
14	27-28	Machine Learning (on Spark): Regression, Classification, Collaborative Filtering, and Clustering.	AR (Apache Spark docs.)
15	29-30	Streaming: Stream Processing – Motivation, Examples, Constraints, and Approaches.	AR
16	31-32	Streaming on Spark: Architecture of Spark Streaming, Stream Processing Model, Example.	AR (Apache Spark docs.)

Select Topics for experiential learning

Topic No.	Select Topics in Syllabus for experiential learning
1	<ul style="list-style-type: none"> Exercises on Distributed Systems – Hadoop; Exercises using Map-reduce model: Map only and reduce only jobs, Standard patterns in

	map reduce models.
2	<ul style="list-style-type: none"> Exercises on NoSQL; Exercises on NoSQL database – Simple CRUD operations and Failure / Consistency tests; Exercises to implement a Web based application that uses NoSQL databases
3	<ul style="list-style-type: none"> Exercises with Pig queries to perform Map-reduce job and understand how to build queries and underlying principles; Exercises on creating Hive databases and operations on Hive, exploring built in functions, partitioning, data analysis
4	<ul style="list-style-type: none"> Exercises on Spark to demonstrate RDD, and operations such as Map, FlatMap, Filter, PairRDD; Typical Spark Programming idioms such as : Selecting Top N, Sorting, and Joins; Exercises on Spark SQL and DataFrames
5	<ul style="list-style-type: none"> Exercises using Spark MLlib: Regression, Classification, Collaborative Filtering, Clustering
6	<ul style="list-style-type: none"> Exercises on Analytics on the Cloud – using AWS, AWS Map-Reduce, AWS data stores / databases.

[**Note:** A few of these topics for experiential learning will be covered by video demonstrations and/or participatory lab sessions operated remotely. Rest of them will be assigned as homework and may be included for evaluation – *see below*. **End of Note.**]

Evaluation Scheme

Legend: EC = Evaluation Component

No	Name	Type	Duration	Weight	Day, Date, Session, Time
EC-1	Assignment I	Take-home, Programming and use of platforms		(10+15) = 25 %	
	Assignment II				
	Quiz I	Online, at scheduled time		5%	
EC-2	Mid-Semester Test	Closed Book	1.5 hours	30%	As per Programme Calendar
EC-3	Comprehensive Exam	Open Book	2.5 hours	40%	As per Programme Calendar

Important Information

Syllabus for Mid-Semester Test (Closed Book): Topics in Weeks 1-8

Syllabus for Comprehensive Exam (Open Book): All topics given in plan of study

Evaluation Guidelines:

1. EC-1 consists of two Assignments and one quiz. Announcements regarding the same will be made in a timely manner.
2. For Closed Book tests: No books or reference material of any kind will be permitted. Laptops/Mobiles of any kind are not allowed. Exchange of any material is not allowed.
3. For Open Book exams: Use of prescribed and reference text books, in original (not photocopies) is permitted. Class notes/slides as reference material in filed or bound form is permitted. All other additional reading materials in filed / bound form are also permitted. However, loose sheets of paper will not be allowed. Use of calculators is permitted in all exams. Laptops/Mobiles of any kind are not allowed. Exchange of any material is not allowed.
4. If a student is unable to appear for the Regular Test/Exam due to genuine exigencies, the student should follow the procedure to apply for the Make-Up Test/Exam. The genuineness of the reason for

absence in the Regular Exam shall be assessed prior to giving permission to appear for the Make-up Exam. Make-Up Test/Exam will be conducted only at selected exam centres on the dates to be announced later.

It shall be the responsibility of the individual student to be regular in maintaining the self-study schedule as given in the course handout, attend the lectures, and take all the prescribed evaluation components such as Assignment/Quiz, Mid-Semester Test and Comprehensive Exam according to the evaluation scheme.